The eXpress Data Path: Programmable Packet Processing for the Linux Kernel

Toke Høiland-Jørgensen
Karlstad University
Sweden
toke.hoiland-jorgensen@kau.se

Jesper Dangaard Brouer
Red Hat
Denmark
brouer@redhat.com

ABSTRACT

Programmable packet processing in software has become increasingly popular, as increases in computational power has made it possible to run custom programmable pipelines at high speeds. Often, this is implemented by so-called kernel bypass techniques, where a userspace application takes complete control of the networking hardware to avoid expensive context switches between kernel and user space. However, this approach makes high-speed packet processing an all-or-nothing proposition, where the processing application has to either handle all network traffic itself, or perform cumbersome re-injection of packets into the kernel.

In this paper, we present an alternative approach to programmable packet processing, where the kernel provides a safe execution environment for custom packet processing applications, executed in the context of the kernel device driver. This system, called the eXpress Data Path (XDP), is part of the mainline Linux kernel, and has been gradually expanded over the last several releases. We show that XDP achieves a maximum packet processing performance of 25 million packets per second on a single core, and illustrate the flexibility of the programming model through three example use cases: inline DDOS protection, layer-3 routing and layer-4 load balancing.

KEYWORDS

XDP, Programmable Networking

1 INTRODUCTION

High-performance packet processing in software has very tight bounds on the time spent processing each packet ($\simeq$ 6 ns per packet at 100 Gbps). Network stacks in general purpose operating systems typically perform too many operations per packet to be able to keep up with this packet rate, which has led to increased popularity of special-purpose toolkits for software packet processing, such as the Data Plane Development Kit (DPDK) [4], which bypass the operating system completely, instead passing control of the network hardware directly to the network application. This approach can significantly improve performance, but has the drawback that it is more difficult to integrate with the existing networking stack, and applications have to re-implement large parts of the stack. In the worst case, this leads to the need to operate two separate stacks; one to perform high-speed packet processing and one to run normal applications, leading to increased management and maintainability costs.

As an alternative to the all-or-nothing kernel bypass packet processing technique, we present a system that adds programmability directly in the operating system networking stack in a cooperative way, making it possible to perform high-speed packet processing that integrates seamlessly with existing systems, while selective leveraging functionality in the operating system. This framework, called the eXpress Data Path (XDP), works by defining a limited execution environment based on an extended version of the Berkeley Packet Filter virtual machine, which allows verified programs to run directly in the device driver context, before the kernel performs any other packet processing tasks.

XDP has been gradually integrated into the Linux kernel over the last several releases. However, no complete architectural description of the system as a whole exists in the literature. In this work we present such a high-level design decription of XDP, its capabilities and how it integrates with the rest of the Linux kernel. Our performance analysis shows raw packet processing performance on a single core of up to 25 million packets per second. In addition, performance scales both downwards, by decreasing CPU usage as the traffic load drops, and upwards, by increasing performance with the number of cores dedicated to packet processing.

Because of its integration with the kernel networking stack, and its high performance, XDP makes it possible to implement applications that previously required their own networking appliance, such as DDOS protection and load balancing, directly on application servers. It also allows a hybrid approach, where certain fast path processing is offloaded to XDP while retaining normal network stack processing for other packets. This allows for high performance processing without sacrificing flexibility. To illustrate these facets of XDP, we supplement our synthetic performance benchmarks with three examples of real-world use cases that can be implemented in XDP: inline DDOS protection on a hypervisor host or application server; full layer-3 packet forwarding capabilities; and layer-4 load balancing by encapsulating packets and resubmitting them to the network.

The rest of this paper is structured as follows: Section 2 first outlines related work. Section 3 then presents the design of XDP and Section 4 presents the raw packet processing performance analysis. Section 5 presents the real-world use cases and their performance. Finally, Section 6 presents current limitations of XDP and future work, and Section 7 concludes.

2 RELATED WORK

3 THE DESIGN OF XDP

XDP is designed to integrate with the Linux networking stack, enhancing it with high-performance programmable hooks in strategic places. This makes it possible to take advantage of the extensive and robust features of the operating system, while adding custom packet processing as required. For this reason, XDP should not be seen as a monolithic system one injects a single program into, but rather a composition of individual parts that operate in concert to achieve the desired outcome.

This section describes the various parts of the XDP system and how they fit together. We begin with a high-level overview of the XDP programming model and how various features of the kernel combine to form a powerful programmable data plane. Following this, we look in detail at the extended BSD Packer Filter (eBPF) virtual machine providing the execution model, and the in-kernel verifier that ensures the safety of loaded eBPF programs. Finally, we give an overview of some of the performance optimisations that are part of the XDP system, and which help ensure high performance.

Figure 1 shows a diagram of how XDP integrates into the Linux kernel, and Figure 2 shows the execution flow of a typical XDP program. Together, they give an overview of the full XDP system, and they will be referenced throughout the exposition below.
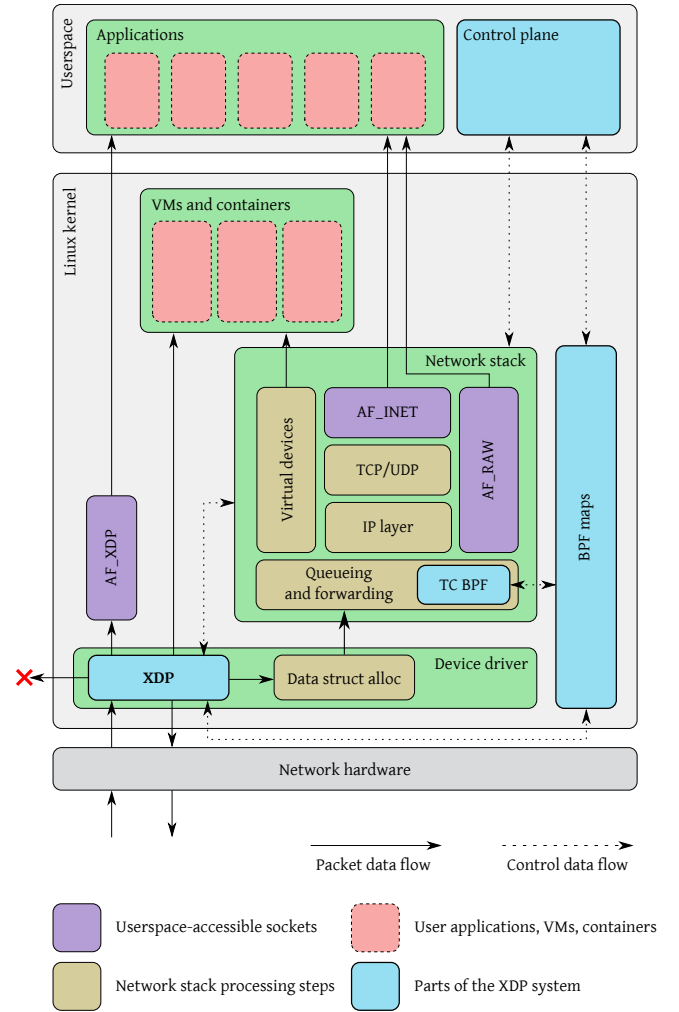

Figure 1: Diagram of how XDP integrates into the Linux kernel network receive path.

3.1 The XDP programming model

The XDP system enables high-performance packet processing integrated tightly with the rest of the Linux networking stack. This makes XDP unique compared to other high-performance software packet processing frameworks, because it makes it possible to selectively leverage features already implemented in Linux, while writing custom programs to perform application-dependent processing, or to accelerate certain parts of the data path. This section gives a conceptual overview of the XDP programming model, explaining how the different parts fit together.

An XDP program is run in the eBPF virtual machine and is entirely event driven. The program is executed directly in context of the device driver, without context switching to userspace. As shown in Figure 1, the program is executed at the earliest possible moment after a packet is received from
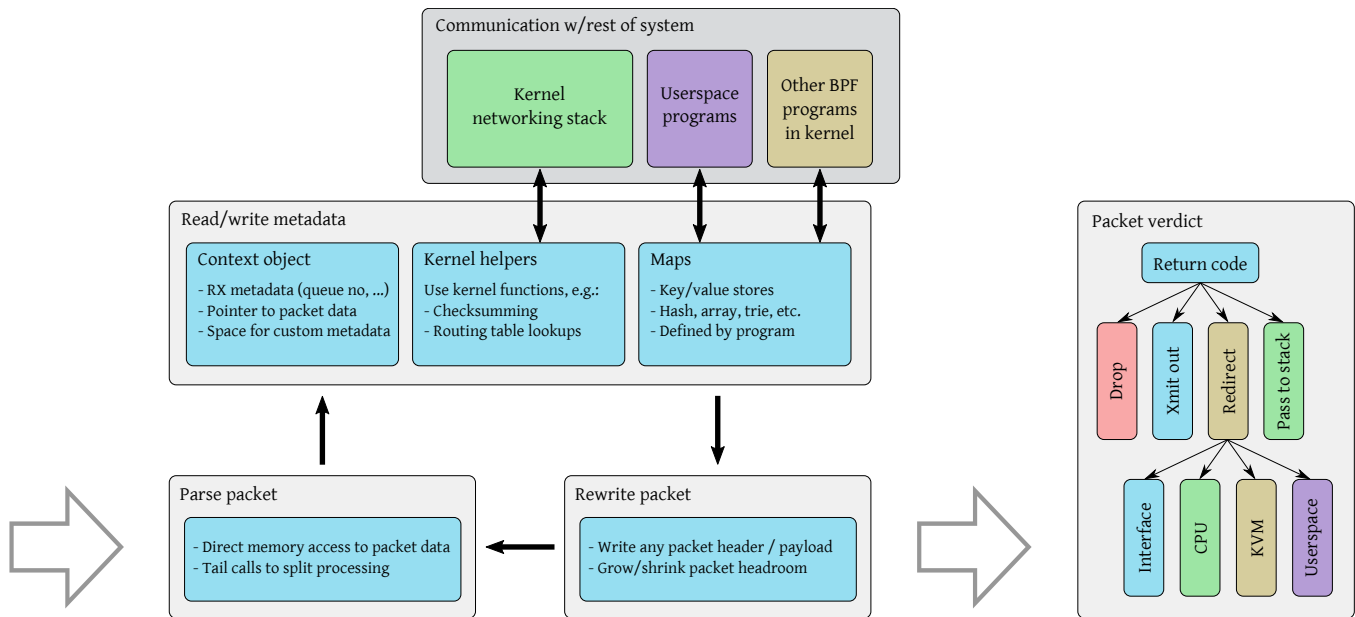
Figure 2: Execution diagram of a typical XDP program. When a packet arrives, the program starts by parsing packet headers to extract the information it will react on. Based on this, combined with information from one or more of the metadata facilities, a packet can be rewritten and a final verdict for the packet determined. The program can alternate between packet parsing, metadata lookup and rewriting, all of which are optional. The final verdict is given in the form of a program return code.

the hardware, before the kernel allocates any data structures or performs any parsing of packet data. This allows for high performance, but the program has to parse raw packet data itself.

Figure 2 shows the various processing steps an XDP program can perform. The program starts its execution with access to a pointer to a context object. This object contains the buffer of the raw packet data, along with metadata fields describing which interface and receive queue the packet came in on, etc. The program typically begins by parsing packet data, and can pass control to a different XDP program, through a so-called *tail call*, which executes the target program and never returns to the caller, thus splitting processing into logical sub-units (based on, say, IP header version). The context object also contains a pointer to a special memory area is available for XDP to store arbitrary metadata that will be available to other XDP programs, as well as to other parts of the kernel and userspace (under certain circumstances). The program can also write any parts of the packet data buffer, including expanding or shrinking the packet to add or remove headers. These three steps (reading, metadata processing, and writing packet data) correspond to the light grey boxes on the left side of Figure 2, and can of course alternate and repeat in a arbitrary ways.

Similar to a regular userspace program, an XDP program ends its execution with a return code, instructing the kernel

what to do with the packet. This is shown on the right-hand side of the figure. There are three simple return codes which either drop the packet, immediately re-transmit it out the same network interface, or allow the packet to be processed by the kernel networking stack. In addition to these simple actions, the XDP program can *redirect* the packet, which controls its further processing. Redirecting can be used to transmit the packet out a different network interface, to pass it to a different CPU to further processing, to pass it to a virtual machine running on the host, or to pass it directly to a special userspace socket without copying. These different ways packets can be passed on are shown with solid lines in Figure 1.

During its execution, and XDP program also has access to kernel facilities that provide helper functions and additional metadata. These allow the XDP program to gather additional metadata through helper functions and maps (dotted lines in Figure 1 and the top part of Figure 2). Helper functions are callbacks implemented in the kernel that an XDP program can call to make use of kernel functionality in its processing. These helpers range serve various purposes, ranging from simple checksum computation and hashing, to full access to the kernel routing table. New helpers are actively been added by the kernel development community in response to user requests, continuously expanding the functionality that XDP programs can make use of.

Maps are key/value stores that are defined by the user before loading an XDP program, and can be referred to from within the eBPF code. Maps are shared, both between different eBPF programs running at various places in the kernel, as well as between eBPF and userspace. The map types include generic hash maps, arrays and radix trees, as well as specialised types containing pointers to eBPF programs, or even recursive pointers to other maps. Maps serve several purposes: they are a persistent data store between invocations of the same eBPF program; a global coordination tool, where eBPF programs in one part of the kernel can update state that changes the behaviour in another; and a communication mechanism between userspace programs and the kernel eBPF programs, similar to the communication between control plane and data plane in other programmable package processing frameworks.

Another piece of the XDP picture is the ability to run eBPF programs in other parts of the kernel. These include packet processing in the Traffic Control (TC) subsystem, where eBPF programs can filter packets after they have been parsed by the kernel, or before they are passed to the hardware from applications. This is marked as "TC BPF" in Figure 1. In addition, eBPF programs can be attached to various places in the kernel that are unrelated to networking (not shown in the figures). These include *cgroups*, which control resource usage for groups of processes (used for implementing containers on Linux, for instance), as well the *tracepoint* and *kprobe* introspection subsystems which allow attaching eBPF programs to arbitrary kernel functions. Because all eBPF programs can share the same set of maps, this makes it possible for XDP programs to react to arbitrary events in the kernel, for instance by dropping packets if processing load increases. Because of this integration, the XDP programming model is considerably more powerful than just the XDP programs itself.

A final important feature of the XDP system is the ability to dynamically load eBPF programs. Because the kernel manages the life cycle of all eBPF programs, they can be dynamically loaded and reloaded at runtime. Combined with dynamic dispatch to other programs using tail calls, this makes it possible to limit the amount of processing actually performed on packets. A processing pipeline can simply split its processing into separate XDP programs and dynamically load and unload them as features are enabled or disabled through control plane configuration. This also makes it possible to dynamically compile programs with hard-coded values derived from configuration, avoiding expensive data structure lookups for common tasks.

The various pieces of the XDP system outlined above combine to form a powerful programmable data plane, with integration into the Linux kernel aiding deployment on existing systems. The following sections describe the eBPF virtual

Table 1: eBPF to x86_64 register mapping.

eBPF	x86_64	eBPF	x86_64	eBPF	x86_64
R0	rax	R4	rcx	R8	r14
R1	rdi	R5	r8	R9	r15
R2	rsi	R6	rbx	R10	rbp
R3	rdx	R7	r13		

machine itself, and the verifier that ensures that loaded programs are safe to run in kernel space.

3.2 The eBPF virtual machine

The eBPF virtual machine is an evolution of the original BSD packet filter (BPF) [6] which has seen extensive use in various packet filtering applications over the last decades. BPF uses a register-based virtual machine to describe filtering actions. This virtual machine has two 32-bit registers and understands 22 different instructions. This makes BPF well-suited for packet filtering operations, but limited as a general purpose virtual machine. eBPF extends the original BPF virtual machine to allow full general purpose execution and efficient just-in-time (JIT) compilation into native machine code. Support for compiling (restricted) C code into eBPF is included in the LLVM compiler suite

The code running in the virtual machine is executed directly in the kernel address space, which makes eBPF useful for a wide variety of tasks in the Linux kernel. The verifier (described in the next section) ensures that user-supplied programs cannot harm the running kernel, which enables a wide array of integrations between the running kernel and the XDP system.

The eBPF modifies the BPF virtual machine as follows:

- The number of registers is increased to eleven, and register widths are increased to 64 bits, with 32-bit sub-registers accessible through certain instructions to provide compatibility with classic BPF programs. The 64-bit registers map one-to-one to hardware registers on all 64-bit architectures supported by the kernel, which eases JIT compilation. For instance, the x86_64 JIT compiler uses the mapping shown in Table 1.
- eBPF adds a *call* instruction for function calls, and adopts the same calling convention as the C language conventions used on the architectures supported by the kernel. Along with the register mapping mentioned above, this makes it possible to map a BPF call instruction to a single native call instruction, enabling function calls to native kernel functions with close to zero overhead. This facility is used by eBPF to support helpers that eBPF programs can call to interact with the kernel while processing.

Table 2: eBPF verifier state variables

Variable	Contains
type	One of the types in Table 3
id	ID for tracking copies of same variable
fixed_offset	Pointer offset (after arithmetic)
range_unsigned	Min and max values (unsigned)
range_signed	Min and max values (signed)
tnum	Mask and value of known bits

The eBPF calling convention is as follows:
- R0 contains the function return value
- R1-R5 contains function arguments
- R6-R9 are callee saved registers that will be preserved across the call
- R10 is a read-only frame pointer to the beginning of the eBPF stack space

A BPF program starts its execution with R1 containing a pointer to a *context* object, the contents of which varies with the type of program. For XDP, this points to a structure that allows the BPF program to access the packet data itself, as well as various items of metadata, including space for arbitrary data that is carried along with the packet and is accessible by other BPF programs that operate on the packet at later stages of processing.

3.3 The eBPF program verifier

As mentioned in the previous section, eBPF code runs directly in the kernel address space, which means that it theoretically has full access to the running kernel and can either crash or compromise this. To avoid this unpleasant situation, the kernel enforces a single entry point for loading all BPF programs (through the bpf() system call). When loading a BPF program it is first analysed by the in-kernel *BPF verifier*, which ensures that the program performs no actions that are unsafe (such as reading arbitrary memory), and that the program will terminate by disallowing loops and limiting the maximum program size. The verifier works by first building a directed acyclic graph (DAG) of the control flow of the program. This DAG is then verified as follows:

First, the verifier performs a depth-first search on the DAG to ensure it contains no loops (no backwards jumps) and that it contains no unsupported or unreachable instructions. Then, in a second pass, the verifier walks all possible paths of the DAG while tracking the state of all registers. The purpose of this second pass is to ensure that the program performs only safe memory accesses, and that any helper functions are called with the right argument types. This is ensured by rejecting programs that perform load or call instructions with invalid arguments. Argument validity is determined by

Table 3: eBPF verifier type annotations. The last column indicates whether pointer arithmetic is allowed for this type of pointer.

Non-pointer types		
Name	Meaning	
NOT_INIT	Not initialised	
SCALAR_VALUE	Any numerical value	
Pointer types		
Name	Pointing to	Arithm
CTX	Context	Yes
MAP	BPF map	No
MAP_VALUE	Value in map	Yes
MAP_VALUE_OR_NULL	Value in map or NULL	No
STACK	Stack frame	Yes
PACKET	Packet data start	Yes
PACKET_END	Packet data end	No

tracking the state of all registers and stack variables through the execution of the program, as explained in the following.

3.3.1 Register state tracking. To track data access, the verifier assigns five state variables to each register, listed in Table 2, with the possible types listed in Table 3. The fixed offset is used to track the result of pointer arithmetic with fixed values, while the ranges and *tnum* are used to track variable offsets of pointers, as well as the ranges of scalar variables.

At the beginning of the program, R1 contains a pointer to the execution context, and its type is a CTX pointer; R10 is a STACK pointer, and all other registers are NOT_INIT. At each execution step, register states are updated based on the operations performed by the program. When a new value is stored to a register, it inherits the state variables of the source of the value. Arithmetic operations on scalar values will affect the value of the *tnum* state variable, which tracks which bits in a register are known, and their value. The *tnum* is a pair of *mask*, which contains the bits whose value is unknown, and a *value* which contains the bits that are known to be set to 1. Load operations set these, for instance loading a byte from memory will result in the top 56 bits being known to be zero, and the bottom 8 bits to be unknown. Arithmetic updates these values according to their operation.

Branches in the instruction tree will update the register state according to the logical operation contained in the branch. For example, a comparison "R1 > 10" compare will set the maximum value of R1 to 10 in one branch, and the minimum value to 11 in the other. If a comparison is performed with a scalar value rather than a constant, the knowledge of which bits are set is used to compute the ranges for the

branches (using the minimum and maximum possible values of unknown bits as appropriate). Finally, a branch that checks whether register with a MAP_VALUE_OR_NULL pointer is different from NULL will turn that register into a pointer with type MAP_VALUE in the *true* branch, which makes it possible to derefence the pointer.

Using the information contained in the state variables, it is possible for the verifier to predict the ranges of memory that it is possible for each load instruction to access. It uses this information to ensure that only safe memory accesses are performed. For pointers to context objects, the execution context of the eBPF program indicates allowed memory offsets for their context objects through a callback performed by the verifier. For map values, the map definition defines the size of the values, which is used to bound the allowed memory accesses. For pointers to stack values, only ranges previously stored on the stack are valid. And finally, for pointers to packet data, only ranges known to be less than the packet length (by appropriate compares against the packet end pointer) are allowed. Any eBPF program that makes memory accesses that the verifier cannot prove are safe are simply rejected at load time. The verifier also uses the range information to enforce aligned memory accesses.

When pointers are copied to other registers, a bounds check on one copy can be used to infer the valid ranges of the other copies, even after the copy occurred. The *id* state variable is used for this purpose for packet access and map value pointers. For packet access, all pointers with the same variable range will have the same *id*, even if their fixed offset differs. Thus, a range check on one copy will mark the same range (minus any differences in fixed offsets) as valid in the other copies. Similarly, for pointers to map values, all copies of a pointer returned from the same map lookup share their *id*, and a check against NULL will be valid for all of them.

3.4 Performance Optimisations in XDP

Much of the performance improvement that XDP represents over the standard Linux networking stack is due to the processing happening before data structures are created and memory allocated. However, there are also a couple of performance enhancement techniques that have specifically been applied over the development of XDP (although some of them also benefit the normal stack). In this section we outline these techniques, and how they apply to XDP.

Packet processing on general-purpose hardware, as is done with XDP, inevitably carries some overhead costs related to getting the packets transferred to the system memory, processing them by the CPU, and sending them back to the hardware. Amortising these costs over multiple packets through bulking is an essential technique to achieve high performance. In Linux, there are two main ways bulking

is achieved: On the receive path, the NAPI mechanism [7] amortises the cost of interrupts from the hardware, by temporarily turning off interrupts each time a packet is received, and instead polling to receive a batch of packets at once. This mechanism has been available in Linux for a long time; but for XDP it carries with it additional benefits, since the XDP program can be executed directly in the NAPI poll context.

A similar issue is seen on the transmit path, where updating the tail pointer in the device ring buffer initiates a transmit operation to the hardware, which carries with it some overhead. To amortise this, XDP uses two mechanisms: When the XDP program indicates that the packet should be transmitted out on the same interface it came in on, it is put into the transmit ring buffer immediately, but the tail pointer update is deferred until the end of the NAPI poll sequence, which causes batching of all packets in the same sequence on transmit as well. When redirecting packets to a different interface, another mechanism is used to achieve the same thing: the redirect mechanism can use a BPF map to lookup the destination interface. When doing so, the map also contains a buffer that will be used to batch packets from subsequent calls, and defer the actual transmission out of the destination interface until the end of the NAPI poll sequence. Using the map structure to achieve this makes it transparent to the calling program and even the device driver.

FIXME: Is the above correct? And is there anything else we need to add to this section?

4 PERFORMANCE EVALUATION

In this section we present our performance evaluation of XDP, using synthetic benchmarks to look at specific aspects of the packet processing capabilities. In the next section, we supplement this with a description and evaluation of a series of real-world use cases.

For all benchmarks, we use a machine equipped with a hexa-core Intel Xeon E5-1650 v4 CPU running at 3.60GHz, and memory modules installed in all four memory slots, to get the full available memory bandwidth. This CPU supports Intel's Data Direct I/O (DDIO) technology, which makes it possible for the networking hardware using Direct Memory Access (DMA) to place data directly into the processor cache, which eliminates cache misses on subsequent processing of the data. The test machine is equipped with two Mellanox ConnectX-5 Ex VPI dual-port 100Gbps network adapters, installed in PCI Express v3 16-lane slots, which ensures that the PCI bus has enough nominal bandwidth available to match the network adapter speed for a single port on each card. However, even so, in some of our tests the performance is bottlenecked at the PCI bus, as we will see below.

We use the TRex packet generator [2] to produce the test traffic. The test machine runs a version of the Linux kernel

that will be released as v4.18. We make details of our setup, links to source code and the raw test data available in an online repository [1].

To show the performance achievable with XDP, we perform the following measurements, which correspond to the four different return codes shown in Figure 2 above:

- Packet drop performance. We install a simple XDP program that drops packets after receiving them, to examine the maximum possible packet processing performance.
- Packet mirroring performance. Here we measure the performance of sending packets out the same network interface that they arrived on.
- Packet forwarding performance. Here we install a simple XDP program that redirects packets out a different interface as they arrive.
- Inline packet processing. We install an XDP program that allows the packets to proceed to the kernel networking stack, to measure the overhead that XDP processing introduces in the kernel receive path.

For all of these tests, we measure the maximum how many packets per second the system can process, as well as the CPU usage required for processing different offered loads. We compare the performance with the `testpmd` example application shipped with DPDK framework, and with the regular Linux kernel network stack (except for the inline processing use case, which is specific to XDP). For all tests, we use minimum-sized (64 bytes) packets, since processing a high number of packets per second is the most challenging, and using larger packets only exacerbates the PCI bus bottleneck.

4.1 Packet Drop Performance

Figure 3 shows the packet processing performance as a function of the number of cores. The baseline performance of XDP for a single core is 26 Mpps, while for DPDK it is 43.5 Mpps. This scales linearly with the number of cores, until it hits the maximum capacity of our traffic generator, at 82.5 Mpps. DPDK reaches this with 3 cores, while for XDP it takes 4. However, we believe it is likely that both XDP and DPDK would be able to continue scaling beyond these limits given a suitable traffic source. **FIXME: Is this true?**

The reason the performance dips at five and six cores is because dividing the traffic over several cores requires using the Receive Side Scaling (RSS) features of the hardware, which can divide up traffic between hardware receive queues (each of which is bound to a separate core) based on packet header information. Using this requires the traffic generator to generate the traffic as separate flows, which means its maximum throughput drops.


Figure 3: Packet drop performance. DPDK uses one core for control tasks, which is why it only goes to 5. The slight downward trend at above 4 cores is because the performance of our packet generator decreases when it has to generate more streams.


Figure 4: Packet forwarding performance.

4.2 Packet Mirroring Performance

TBD. FIXME: Can we do this with DPDK?

4.3 Packet Forwarding Performance

Figure 4 shows the packet forwarding performance. Again, we see an almost linear scaling with the first couple of cores, tapering off to sub-linear performance increases as more cores are added. We attribute this to **FIXME**.

5 REAL-WORLD USE CASES

In this section we describe and evaluate three real-world use cases that showcase various ways in which XDP can used to implement useful applications or features. These use cases have all seen real-world deployment in one form or another, although we use simplified versions in our evaluation so as to be able to make the code available.

The three use cases are inline Denial of Service (DoS) mitigation, a software layer-3 router, and an application load-balancer, and are described in turn in each of the following subsections.

5.1 Inline DoS Mitigation

DoS attacks continue to plague the internet, typically in the form of distributed attacks (DDoS attacks) from compromised devices attached to the internet. There are various ways to mitigate them, but having some kind of service that protects critical infrastructure is typically at least a component of DDoS mitigation **FIXME: citation needed**. This can be done by filtering all traffic through a device that distinguish legitimate traffic from attacks, and drop the attack packets before they reach the application server.

The obvious problem with protecting against a DDoS attack is the sheer scale of the traffic that needs to be dropped, which leads to the use of expensive appliances to handle the dropping. However, with XDP we have another option: installing the traffic filter directly on the application server in the form of an XDP program, which is possible without any other modifications to the server. In the case of a virtual machine deployment, the filter can even be installed on the hypervisor, and thus protect all virtual machines running on the host.

How to identify attack traffic is not our focus here, but rather the mechanism that drops the unwanted traffic. We simply assume that there exists *some* way of identifying which traffic is legitimate and which isn't. Because of the dynamic nature of XDP, there are many options available: BPF maps could be used for while-listing or black-listing, the packet payload could be inspected, etc. It is even possible to load new XDP programs especially tailored towards dropping a particular kind of attack traffic as the need arises. However, in our example, we simply assume that a server hosts a TCP-based service, and so consider all UDP traffic as unwanted.

To test the performance of such a solution, we use the Netperf benchmarking tool [5]. This is equipped with a TCP-based round-trip benchmark, which opens a TCP connection and sends a small payload which is echoed back from the server, repeating as soon as a reply is received. The output is the number of transactions per second, which is a good proxy for an interactive TCP session, such as small remote procedure calls.

We first measure the baseline performance of this test, and see that we are able to perform around 25.000 transactions per second with no competing traffic. We then, from a separate machine, offer an increasing load of small UDP packets, which simulates the DoS attack, and measure how the TCP application is affected. The network interface is configured to redirect all traffic to a single receive queue, so all packets


Figure 5: DDOS performance. Number of TCP transactions per second with different levels of UDP packet floods being sent to the server.

are processed by the same CPU core. This is important in this use case, because the machine will be running other applications on the remaining cores, and so we don't want to dedicate too many resources to just packet processing[1].

The results of this is shown in Figure 5. It is clearly seen that without the XDP filter, performance drops significantly, being halved at 2.5 Mpps and effectively zero at just below 3 Mpps. However, with the XDP filter in place, performance is kept above 20.000 transactions per second up to 19 Mpps, after which it drops rapidly. The small increase in baseline performance (from 24,500 to 25,500 transactions per second) is because the XDP program uses the *redirect* feature to move the network stack processing of the application traffic to a separate CPU core, to make sure the application performance is not impacted by handling of the DoS traffic.

As these results show, it is quite feasible to perform this kind of filtering in XDP, and comfortably handle packet rates above 10 Gbps of DoS traffic (with minimum packet sizes) on a single CPU core. Deploying DoS mitigation this way leads to increased flexibility, both because of the programmability offered by XDP, but also because it eliminates the need for at separate appliance to scrub the traffic.

5.2 Software Router

The second use case is that of a software router. The Linux kernel already contains a full-featured routing table, which includes support for policy routing, source-specific routing, equal-cost multipath load balancing, and more. And routing daemons such as Bird or FRR implement a variety of routing control plane protocols, which makes it quite feasible to turn a Linux system into a full-featured software router. However,

---

[1]The actual number of cores dedicated to packet processing will of course have to depend on the volume of traffic the filter should be able to process.

thus far the data plane forwarding performance has kept this from being feasible at very high rates.

Because of the rich ecosystem for routing on Linux, improving performance of the kernel data plane is desirable, as re-implementing the routing stack in another data processing framework carries a high cost. Thus, XDP is particularly suited for this task. This is also the reason why one of the first kernel helpers introduced to XDP is a routing table lookup function. This function makes it possible to perform full routing table lookups directly from XDP. The result of the lookup is an egress interface and a next-hop MAC address, which makes it possible for the XDP program to immediately forward the packet if the lookup succeeds. If no next-hop MAC is known (because neighbour lookup hasn't been performed yet), the XDP program can instead pass the packet to the networking stack, which will perform the neighbour lookup, allowing subsequent packets to be forwarded by the XDP fast path.

To show the performance of the software router use case, we use the XDP routing example that is included in the Linux kernel source and compare its performance to the native Linux networking stack routing. The example simply parses the IP header of an incoming packet, does a routing table lookup, and immediately forwards the packet if a match is found, or passes it up to the network stack if not. We perform two tests: one with a single route installed in the routing table and all packets set to the same destination address. And another where we import a full dump of the global BGP routing table from routeviews.org, where we set all next-hop addresses to the same value, but vary the source addresses of the packets being routed. In the first test we compare the performance with the layer-3 forwarding example included with the DPDK source code; but because this example only supports a fixed routing table, we do not perform the full BGP table comparison with DPDK.

The performance of this use case is seen in Figure **FIXME**.

5.3   Load Balancer

For the load balancer use case, we use the XDP component of the Katran load balancer [3] recently released as open source by Facebook. This works by announcing a virtual IP address to the world, which is then routed to the load balancer. At the load balancer, the XDP program hashes the layer-4 packet header using a stable hashing algorithm, which becomes a destination application server. The packet is then encapsulated in a new IP header and sent to the application server, which is responsible for decapsulating it, processing the request, and replying directly to the originator of the request. This means that the XDP program can perform the packet encapsulation, and immediately return the packet out the same interface, as in the performance test in Section 4.2.

To test this use case, we configure the Katran XDP program with a fixed number of destination hosts, and run it on our test machine. Once the packets are returned to the network they are simply dropped; we are only interested in the performance of the load balancer component itself. This performance is shown in Figure **FIXME**.

6   LIMITATIONS AND FUTURE WORK

As we have shown above, XDP has excellent performance, and is already useful in a variety of real-world use cases. However, XDP is also relatively new technology that is being incrementally added to a general purpose operating system kernel. As such, it is constantly evolving, and still has some rough edges. In this section we present some of the known limitations, and future directions for XDP's development. These include: User experience and debugging; driver support; improvements to NAPI; and better handling of rate transitions. Each of the following subsections discuss each of these issues in turn.

6.1   User Experience and Debugging

6.2   Driver Support

6.3   Improvements to NAPI

6.4   Handling of Rate Transitions

Currently, XDP offers no mechanism to gracefully handle rate transitions, i.e., cases where packets are forwarded to a destination that cannot handle the rate of incoming packets. This can happen if, e.g., a 100 Gbps ingress interface attempts to forward packets to a slower 10 Gbps interface, or to a virtual interface or userspace application that cannot process them fast enough.

Since the XDP program signals the packet destination via its return code, execution has already ended when the packet forwarding is attempted, and so a failure due to a full hardware buffer or a busy userspace application cannot be gracefully handled, and will simply result in the packet being silently dropped. This is obviously not ideal, and exploring ways to gracefully deal with this is planned for the future.

7   CONCLUSIONS

REFERENCES

[1] A. Author. 2018. Online repository of test scripts. URL omitted for blinded review.
[2] Cisco. 2018. TRex Traffic Generator. https://trex-tgn.cisco.com/
[3] Facebook. 2018. Katran source code repository. https://github.com/facebookincubator/katran
[4] Linux Foundation. 2018. Data Plane Development Kit. https://dpdk.org/
[5] Rick Jones. 2018. Netperf. Open source benchmarking software. http://www.netperf.org/
[6] Steven McCanne and Van Jacobson. 1993. The BSD Packet Filter: A New Architecture for User-level Packet Capture.. In *USENIX winter*, Vol. 93.

[7] Jamal Hadi Salim, Robert Olsson, and Alexey Kuznetsov. 2001. Beyond Softnet. In *Proceedings of the 5th Annual Linux Showcase & Conference*. USENIX. https://wiki.linuxfoundation.org/networking/napi/