The eXpress Data Path: Programmable Packet Processing for the Linux Kernel

Toke Høiland-Jørgensen
Karlstad University
Sweden
toke.hoiland-jorgensen@kau.se

Jesper Dangaard Brouer
Red Hat
Denmark
brouer@redhat.com

ABSTRACT

Programmable packet processing in software has become increasingly popular, as increases in computational power has made it possible to run custom programmable pipelines at high speeds. This is generally implemented by so-called kernel bypass techniques, where a userspace application takes complete control of the networking hardware to avoid expensive context switches between kernel and user space.

However, this approach makes high-speed packet processing an all-or-nothing proposition, where the processing application has to handle all network traffic, reinventing a lot of functionality already present in the kernel. Interoperability with other applications is also difficult, requiring cumbersome re-injection of packets into the kernel.

In this paper, we present an alternative approach to programmable packet processing, where the kernel provides a safe execution environment for custom packet processing applications, executed directly in the context of the kernel device driver. This system, called the eXpress Data Path (XDP), is part of the mainline Linux kernel, and has been gradually expanded over the last several releases. We describe the design of the system and how it integrates with the rest of the kernel. Through a range of real-world examples we illustrate the flexibility of this model and how it can be used to implement a wide range of applications with performance measure in tens of millions of packets per second.

KEYWORDS

XDP, Programmable Networking

1 INTRODUCTION

High-performance packet processing in software has very tight bounds on the time spent processing each packet ($\simeq$ 67 ns per packet at 10 Gbps). Network stacks in general purpose operating systems typically perform way too many operations per packet to be able to keep up with this packet rate, which has led to the introduction of special-purpose networking toolkits for software packet processing, such as DPDK and Netmap. However, these toolkits have the drawback that they are difficult to integrate with the existing networking stack, leading to the need to re-implement large parts of the stack.

We present an alternative to previous approaches: A novel way to integrate programmable packet processing directly into the networking stack in a cooperative way, making it possible to perform high-speed packet processing that integrates seamlessly with existing applications. This framework, called the eXpress Data Path (XDP), works by defining a limited execution environment based on an extended version of the Berkeley Packet Filter bytecode language, which allows verified programs to run directly in kernel context before the normal packet processing in the networking stack.

This makes it possible to implement applications that previously required their own appliance, such as DDOS protection and load balancing, directly on application servers. It also allows a hybrid approach, where certain fast path processing is offloaded to XDP while retaining normal network stack processing for other packets. This allows for exceptionally high throughput and low latency processing without sacrificing flexibility.

We present the design of XDP and its capabilities and integration with the Linux kernel. We then present a performance evaluation that consists of micro-benchmarks showing packet processing scaling beyond 20 Mpps on a single core as well as two real-world use cases: inline DDOS protection and layer-3 packet forwarding.

2 RELATED WORK

3 THE DESIGN OF XDP

XDP is designed to integrate with the Linux networking stack, enhancing it with high-performance programmable hooks in strategic places. This makes it possible to take advantage of the extensive and robust features of the operating system, while adding custom packet processing as required. For this reason, XDP should not be seen as a monolithic system one injects a single program into, but rather a composition of individual parts that operate in concert to achieve the desired outcome.

This section describes the various parts of the XDP system and how they fit together. We begin with a high-level overview of the XDP programming model and how various features of the kernel combine to form a powerful programmable data plane. Following this, we look in detail at the extended BSD Packer Filter (eBPF) virtual machine providing the execution model, and the in-kernel verifier that ensures the safety of loaded eBPF programs.

3.1 The XDP programming model

The XDP system enables high-performance packet processing integrated tightly with the rest of the Linux networking stack. This makes XDP unique compared to other high-performance software packet processing frameworks, because it makes it possible to selectively leverage features already implemented in Linux, while writing custom programs to perform application-dependent processing, or to accelerate certain parts of the data path. This section gives a conceptual overview of the XDP programming model, explaining how the different parts fit together.

An XDP program is run in the eBPF virtual machine and is entirely event driven. The program is executed directly in context of the device driver, without context switching to userspace. The program is executed at the earliest possible moment after a packet is received from the hardware, before the kernel allocates any data structures or performs any parsing of packet data. This allows for high performance, but the program has to parse raw packet data itself.

Figure 1 gives a conceptual overview of the execution of a typical XDP program. The program starts its execution with access to a pointer to a context object. This object contains the buffer of the raw packet data, along with metadata fields describing which interface and receive queue the packet came in on, etc. The program can read and write any parts of the packet data buffer, including expanding or shrinking the packet to add or remove headers. It can also pass control to a different XDP program, through a so-called *tail call*, which executes the target program and never returns to the caller, thus splitting processing into logical sub-units (based on, say, IP header version). The context object also contains

a pointer to a special memory area is available for XDP to store arbitrary metadata that will be available to other XDP programs, as well as to other parts of the kernel and userspace (under certain circumstances).

Similar to a regular userspace program, an XDP program ends its execution with a return code, instructing the kernel what to do with the packet. There are three simple return codes which either drop the packet, immediately re-transmit it out the same network interface, or allow the packet to be processed by the kernel networking stack. In addition to these simple actions, the XDP program can *redirect* the packet, which controls its further processing. Redirecting can be used to transmit the packet out a different network interface, to pass it to a different CPU to further processing, to pass it to a virtual machine running on the host, or to pass it directly to a special userspace socket without copying.

During its execution, and XDP program also has access to kernel facilities that provide helper functions and additional metadata. Figure 2 shows a diagram of how XDP integrates into the full system. The XDP hook in the driver is the first thing that is executed after a packet is received from the network hardware. This program can control how further packet processing is performed as described above (solid lines). In addition to this, the facilities given to interact with the kernel allow the XDP program to gather additional metadata (dotted lines), through helper functions and maps. Helper functions are callbacks implemented in the kernel that an XDP program can call to make use of kernel functionality in its processing. These helpers range serve various purposes, ranging from simple checksum computation and hashing, to full access to the kernel routing table. New helpers are actively been added by the kernel development community in response to user requests, continuously expanding the functionality that XDP programs don't have to re-implement themselves.

Maps are key/value stores that are defined by the user before loading an XDP program, and can be referred to from within the eBPF code. Maps are shared, both between different eBPF programs running at various places in the kernel, as well as between eBPF and userspace. The map types include generic hash maps, arrays and radix trees, as well as specialised types containing pointers to eBPF programs, or even recursive pointers to other maps. Maps serve several purposes: they are a persistent data store between invocations of the same eBPF program; a global coordination tool, where eBPF programs in one part of the kernel can update state that changes the behaviour in another; and a communication mechanism between userspace programs and the kernel eBPF programs, similar to the communication between control plane and data plane in other programmable package processing frameworks. These different communication paths are marked with dotted lines in Figure 2.
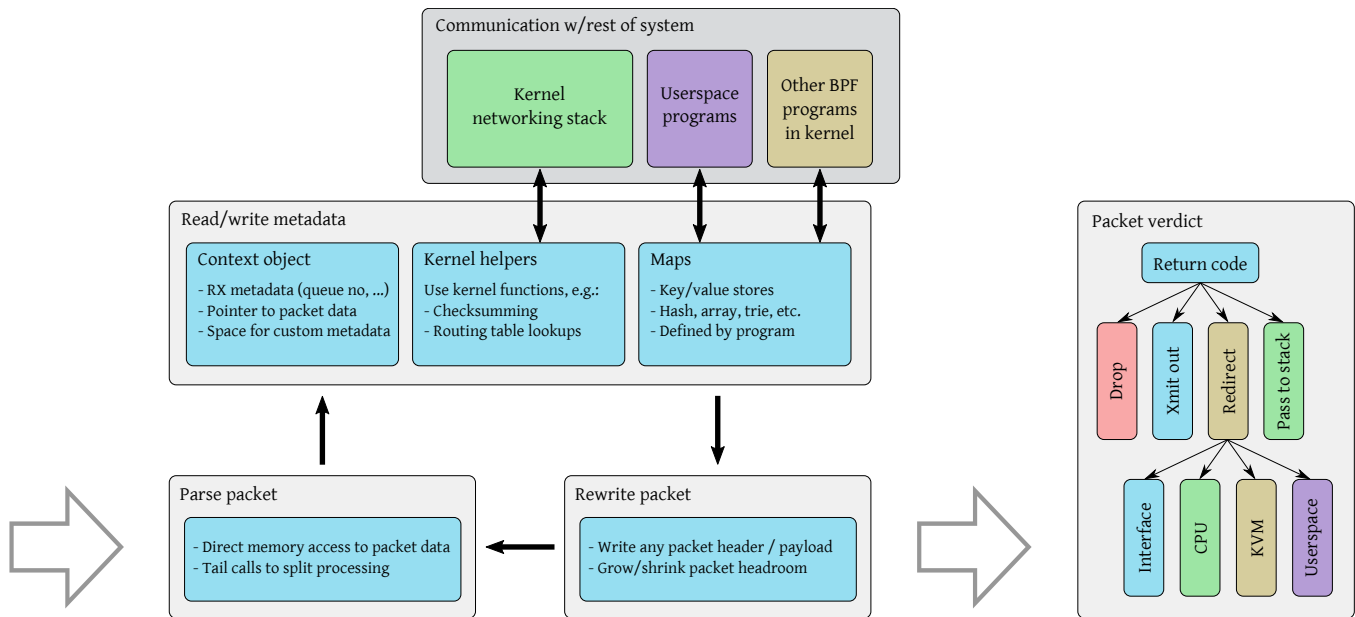
Figure 1: Execution diagram of a typical XDP program. When a packet arrives, the program starts by parsing packet headers to extract the information it will react on. Based on this, combined with information from one or more of the metadata facilities, a packet can be rewritten and a final verdict for the packet determined. The program can alternate between packet parsing, metadata lookup and rewriting, all of which are optional. The final verdict is given in the form of a program return code.

Another piece of the XDP picture is the ability to run eBPF programs in other parts of the kernel. These include packet processing in the Traffic Control (TC) subsystem, where eBPF programs can filter packets after they have been parsed by the kernel, or before they are passed to the hardware from applications. This is marked as "TC BPF" in the figure. In addition, eBPF programs can be attached to various places in the kernel that are unrelated to networking (not shown on the figure). These include *cgroups*, which control resource usage for groups of processes (used for implementing containers on Linux, for instance), as well the *tracepoint* and *kprobe* introspection subsystems which allow attaching eBPF programs to arbitrary kernel functions. Because all eBPF programs can share the same set of maps, this makes it possible for XDP programs to react to arbitrary events in the kernel, for instance by dropping packets if processing load increases. Because of this integration, the XDP programming model is considerably more powerful than just the XDP programs itself.

A final important feature of the XDP system is the ability to dynamically load eBPF programs. Because the kernel manages the life cycle of all eBPF programs, they can be dynamically loaded and reloaded at runtime. Combined with dynamic dispatch to other programs using tail calls, this makes it possible to limit the amount of processing actually performed on packets. A processing pipeline can simply split

its processing into separate XDP programs and dynamically load and unload them as features are enabled or disabled through control plane configuration. This also makes it possible to dynamically compile programs with hard-coded values derived from configuration, avoiding expensive data structure lookups for common tasks.

The various pieces of the XDP system outlined above combine to form a powerful programmable data plane, with integration into the Linux kernel aiding deployment on existing systems. The following sections describe the eBPF virtual machine itself, and the verifier that ensures that loaded programs are safe to run in kernel space.

3.2 The eBPF virtual machine

The eBPF virtual machine is an evolution of the original BSD packet filter (BPF) [1] which has seen extensive use in various packet filtering applications over the last decades. BPF uses a register-based virtual machine to describe filtering actions. This virtual machine has two 32-bit registers and understands 22 different instructions. This makes BPF well-suited for packet filtering operations, but limited as a general purpose virtual machine. eBPF extends the original BPF virtual machine to allow full general purpose execution and efficient just-in-time (JIT) compilation into native machine code. Support for compiling (restricted) C code into eBPF is included in the LLVM compiler suite
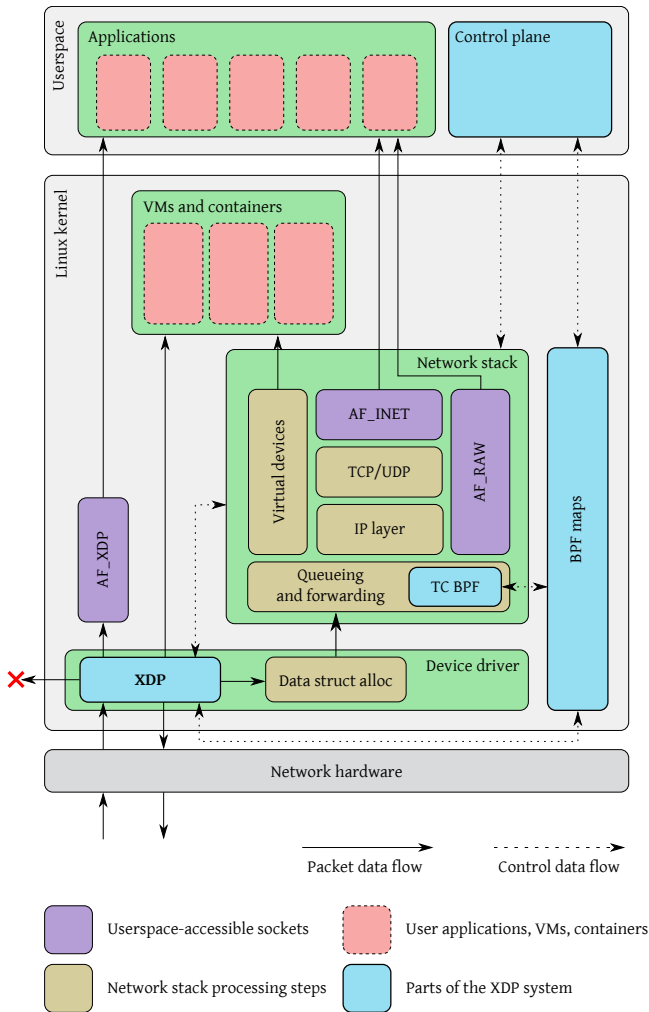
Toke Høiland-Jørgensen and Jesper Dangaard Brouer

Figure 2: Diagram of how XDP integrates into the Linux kernel network receive path.

The code running in the virtual machine is executed directly in the kernel address space, which makes eBPF useful for a wide variety of tasks in the Linux kernel. The verifier (described in the next section) ensures that user-supplied programs cannot harm the running kernel, which enables a wide array of integrations between the running kernel and the XDP system.

The eBPF modifies the BPF virtual machine as follows:

- The number of registers is increased to eleven, and register widths are increased to 64 bits, with 32-bit sub-registers accessible through certain instructions to provide compatibility with classic BPF programs. The 64-bit registers map one-to-one to hardware registers on all 64-bit architectures supported by the kernel, which eases JIT compilation. For instance, the x86_64 JIT compiler uses the mapping shown in Table 1.

Table 1: eBPF to x86_64 register mapping.

eBPF	x86_64
R0	rax
R1	rdi
R2	rsi
R3	rdx
R4	rcx
R5	r8
R6	rbx
R7	r13
R8	r14
R9	r15
R10	rbp

- eBPF adds a *call* instruction for function calls, and adopts the same calling convention as the C language conventions used on the architectures supported by the kernel. Along with the register mapping mentioned above, this makes it possible to map a BPF call instruction to a single native call instruction, enabling function calls to native kernel functions with close to zero overhead. This facility is used by eBPF to support helpers that eBPF programs can call to interact with the kernel while processing.
  The eBPF calling convention is as follows:
  – R0 contains the function return value
  – R1-R5 contains function arguments
  – R6-R9 are callee saved registers that will be preserved across the call
  – R10 is a read-only frame pointer to the beginning of the eBPF stack space

A BPF program starts its execution with R1 containing a pointer to a *context* object, the contents of which varies with the type of program. For XDP, this points to a structure that allows the BPF program to access the packet data itself, as well as various items of metadata, including space for arbitrary data that is carried along with the packet and is accessible by other BPF programs that operate on the packet at later stages of processing.

3.3 The eBPF program verifier

As mentioned in the previous section, eBPF code runs directly in the kernel address space, which means that it theoretically has full access to the running kernel and can either crash or compromise this. To avoid this unpleasant situation, the kernel enforces a single entry point for loading all BPF programs (through the bpf() system call). When loading a BPF program it is first analysed by the in-kernel *BPF verifier*, which ensures that the program performs no actions that

are unsafe (such as reading arbitrary memory), and that the program will terminate by disallowing loops and limiting the maximum program size. The verifier works by first building a directed acyclic graph (DAG) of the control flow of the program. This DAG is then verified as follows:

First, the verifier performs a depth-first search on the DAG to ensure it contains no loops (no backwards jumps) and that it contains no unsupported or unreachable instructions. Then, in a second pass, the verifier walks all possible paths of the DAG while tracking the state of all registers. The purpose of this second pass is to ensure that the program performs only safe memory accesses, and that any helper functions are called with the right argument types. This is ensured by rejecting programs that perform load or call instructions with invalid arguments. Argument validity is determined by tracking the state of all registers and stack variables through the execution of the program, as explained in the following.

3.3.1 Register state tracking. To track data access, the verifier assigns five state variables to each register, listed in Table 2, with the possible types listed in Table 3. The fixed offset is used to track the result of pointer arithmetic with fixed values, while the ranges and *tnum* are used to track variable offsets of pointers, as well as the ranges of scalar variables.

At the beginning of the program, R1 contains a pointer to the execution context, and its type is a CTX pointer; R10 is a STACK pointer, and all other registers are NOT_INIT. At each execution step, register states are updated based on the operations performed by the program. When a new value is stored to a register, it inherits the state variables of the source of the value. Arithmetic operations on scalar values will affect the value of the *tnum* state variable, which tracks which bits in a register are known, and their value. The *tnum* is a pair of *mask*, which contains the bits whose value is unknown, and a *value* which contains the bits that are known to be set to 1. Load operations set these, for instance loading a byte from memory will result in the top 56 bits being known to be zero, and the bottom 8 bits to be unknown. Arithmetic updates these values according to their operation.

Branches in the instruction tree will update the register state according to the logical operation contained in the branch. For example, a comparison "R1 > 10" compare will set the maximum value of R1 to 10 in one branch, and the minimum value to 11 in the other. If a comparison is performed with a scalar value rather than a constant, the knowledge of which bits are set is used to compute the ranges for the branches (using the minimum and maximum possible values of unknown bits as appropriate). Finally, a branch that checks whether register with a MAP_VALUE_OR_NULL pointer is different from NULL will turn that register into a pointer with

type MAP_VALUE in the *true* branch, which makes it possible to derefence the pointer.

Using the information contained in the state variables, it is possible for the verifier to predict the ranges of memory that it is possible for each load instruction to access. It uses this information to ensure that only safe memory accesses are performed. For pointers to context objects, the execution context of the eBPF program indicates allowed memory offsets for their context objects through a callback performed by the verifier. For map values, the map definition defines the size of the values, which is used to bound the allowed memory accesses. For pointers to stack values, only ranges previously stored on the stack are valid. And finally, for pointers to packet data, only ranges known to be less than the packet length (by appropriate compares against the packet end pointer) are allowed. Any eBPF program that makes memory accesses that the verifier cannot prove are safe are simply rejected at load time. The verifier also uses the range information to enforce aligned memory accesses.

When pointers are copied to other registers, a bounds check on one copy can be used to infer the valid ranges of the other copies, even after the copy occurred. The *id* state variable is used for this purpose for packet access and map value pointers. For packet access, all pointers with the same variable range will have the same *id*, even if their fixed offset differs. Thus, a range check on one copy will mark the same range (minus any differences in fixed offsets) as valid in the other copies. Similarly, for pointers to map values, all copies of a pointer returned from the same map lookup share their *id*, and a check against NULL will be valid for all of them.

Table 2: eBPF verifier state variables

Variable	Contains
type	One of the types in Table 3
id	ID for tracking copies of same variable
fixed_offset	Pointer offset (after arithmetic)
range_unsigned	Min and max values (unsigned)
range_signed	Min and max values (signed)
tnum	Mask and value of known bits

4 PERFORMANCE EVALUATION

4.1 Micro-benchmarks

4.2 Comparison with DPDK/netmap

5 REAL-WORLD USE CASES

5.1 DDOS mitigation

5.2 Packet forwarding layer 2/3

- Helper functions into bridging / routing code

Table 3: eBPF verifier type annotations. The last column indicates whether pointer arithmetic is allowed for this type of pointer.

Non-pointer types		
Name	Meaning	
NOT_INIT	Not initialised	
SCALAR_VALUE	Any numerical value	
Pointer types		
Name	Pointing to	Arithm
CTX	Context	Yes
MAP	BPF map	No
MAP_VALUE	Value in map	Yes
MAP_VALUE_OR_NULL	Value in map or NULL	No
STACK	Stack frame	Yes
PACKET	Packet data start	Yes
PACKET_END	Packet data end	No

- Layer 2 also useful for VMs

5.3 Load-balancer

- XDP_TX
- XDP_REDIRECT to CPU/VM

6 CONCLUSIONS

REFERENCES

[1] Steven McCanne and Van Jacobson. 1993. The BSD Packet Filter: A New Architecture for User-level Packet Capture.. In *USENIX winter*, Vol. 93.